

# Miruna Oprescu

---

PhD Candidate in Computer Science

Cornell University, Cornell Tech

✉ [miruna@cs.cornell.edu](mailto:miruna@cs.cornell.edu) | [🌐 mirunaoprescu.com](https://mirunaoprescu.com)

## SUMMARY

I develop robust and reliable machine learning algorithms for causal inference and data-driven decision-making. My research tackles challenging settings, such as experiments with imperfect compliance, sensitivity and risk analysis in observational data, reinforcement learning for dynamic decision-making, and datasets with complex structure. My work aims to create trustworthy and effective models for real-world, data-limited, and high-impact settings.

## EDUCATION

### Cornell University, Cornell Tech

Fall 2021 – Present

Ph.D. Candidate in Computer Science. GPA: 4.00

Department of Energy Computational Science Graduate Fellow

M.S. in Computer Science, December 2024 (awarded en route to Ph.D.)

### Harvard University

May 2015

Joint A.B. degree in Physics and Mathematics. Minor in Computer Science

*Cum laude in field* with High Honors in Physics and Mathematics

## EXPERIENCE

### Cornell University, Cornell Tech

New York, NY

Graduate Research Assistant

Fall 2021 – Present

- Research in causal inference, machine learning, and robust data-driven decision-making. Adviser: Nathan Kallus.

### Netflix

New York City, NY

Machine Learning Research Intern

Summer 2025

- Building causal models of user engagement and retention, with applications to content value and membership dynamics. Mentors: Nathan Kallus, David Hubbard.

### Brookhaven National Laboratory

Brookhaven, NY

Research Intern

Summer 2024

- Developed causal inference methods for spatio-temporal applications in Earth Science. Mentor: Shinjae Yoo.

### Netflix

Los Gatos, CA

Machine Learning Research Intern

Summer 2022

- Developed and built a causal machine learning model for quantifying the causal effect of watching a new title on long term user engagement. Mentors: Sudeep Das, Aish Fenton.

### Microsoft Research

Cambridge, MA

Senior Data and Applied Scientist

2019 – 2021

Data and Applied Scientist II

2017 – 2019

- Conducted research on machine learning-based causal inference techniques, contributing to top conference publications.
- Developed and published causal inference algorithms as a core contributor to the [EconML](#) library, supporting high-impact projects across various industries.
- Researched and improved subseasonal weather forecasting models, with results published in leading journals and conferences.

**Microsoft**

Software Development Engineer

Cambridge, MA

2015 – 2017

- Developed and published scalable machine learning algorithms as a core contributor to [MMLSpark](#), the Microsoft Machine Learning Library for Apache Spark.

**Johns Hopkins University**

Research Intern

Baltimore, MD

Summer 2014

- Developed clustering algorithms for clinical time series data to predict septic shock and created a dynamic web application for visualizing clusters and analyzing health data. Mentor: Suchi Saria.

**SELECTED  
PUBLICATIONS**

† - equal contribution, ‡ - alphabetical authors

- [1] **Miruna Oprescu** and Nathan Kallus. Estimating heterogeneous treatment effects by combining weak instruments and observational data. *Advances in Neural Information Processing Systems*, 37:118777–118806, 2025.
- [2] Andrew Bennett<sup>†</sup>, Nathan Kallus<sup>†</sup>, **Miruna Oprescu**<sup>†</sup>, Wen Sun<sup>†</sup>, and Kaiwen Wang<sup>†</sup>. Efficient and sharp off-policy evaluation in robust markov decision processes. *Advances in Neural Information Processing Systems*, 37:112962–113000, 2025.
- [3] **Miruna Oprescu**, Jacob Dorn, Marah Ghoummaid, Andrew Jesson, Nathan Kallus, and Uri Shalit. B-learner: Quasi-oracle bounds on heterogeneous causal effects under hidden confounding. In *Proceedings of the 40th International Conference on Machine Learning*, pages 26599–26618. PMLR, 2023.
- [4] Nathan Kallus<sup>†</sup> and **Miruna Oprescu**<sup>†</sup>. Robust and agnostic learning of conditional distributional treatment effects. In *International Conference on Artificial Intelligence and Statistics*, pages 6037–6060. PMLR, 2023.
- [5] Keith Battocchi<sup>†</sup>, Eleanor Dillon<sup>†</sup>, Maggie Hei<sup>†</sup>, Greg Lewis<sup>†</sup>, **Miruna Oprescu**<sup>†</sup>, and Vasilis Syrgkanis<sup>†</sup>. Estimating the long-term effects of novel treatments. *Advances in Neural Information Processing Systems*, 34:2925–2935, 2021.
- [6] **Miruna Oprescu**<sup>†</sup>, Vasilis Syrgkanis<sup>†</sup>, and Zhiwei Steven Wu<sup>†</sup>. Orthogonal random forest for causal inference. In *International Conference on Machine Learning*, pages 4932–4941. PMLR, 2019.
- [7] Vasilis Syrgkanis, Victor Lei, **Miruna Oprescu**, Maggie Hei, Keith Battocchi, and Greg Lewis. Machine learning estimation of heterogeneous treatment effects with instruments. In *Advances in Neural Information Processing Systems*, pages 15193–15202, 2019. **Spotlight presentation.**
- [8] **Miruna Oprescu**<sup>†</sup>, Vasilis Syrgkanis<sup>†</sup>, Keith Battocchi<sup>†</sup>, Maggie Hei<sup>†</sup>, and Greg Lewis<sup>†</sup>. EconML: A Machine Learning Library for Estimating Heterogeneous Treatment Effects. In *CausalML Workshop, NeurIPS*, 2019. **Spotlight presentation.**

**TALKS***Causal Inference for Spatiotemporal Interventions*

SIAM Conference on Computational Science and Engineering (CSE25), 2025.

*Reliable Treatment Effect Estimation Using Weak Instruments and Observational Data*  
Workshop in Operations Research and Data Science (WORDS), Duke University, 2024.*Uncertainty Quantification in Causal Inference: Sharp and Efficient Bounds on Heterogeneous Causal Effects Under Hidden Confounding*

Computational Science Seminar, Brookhaven National Laboratory, 2023.

*Causal Inference and Machine Learning in Practice with EconML and CausalML: Industrial Use Cases at Microsoft, TripAdvisor, Uber*  
The SIGKDD Conference on Knowledge Discovery & Data Mining, 2021.

*EconML: A Machine Learning Library for Estimating Heterogeneous Treatment Effects*  
Open Data Science Conference East, 2019.

*MMLSpark: Lessons from Building a SparkML Compatible Machine Learning Library*  
Spark Summit Europe, 2017

<b>HONORS &amp; AWARDS</b>	Department of Energy Computational Science Graduate Fellowship	2022 – 2026
	Meta PhD Research Fellowship Finalist	2022
	<i>cum laude</i> , Harvard University	2015
	High Honors, Harvard University Physics Department	2015
	Derek C. Bok Award for Distinction in Teaching ( <i>Data Science</i> ), Harvard	2014

<b>SERVICE</b>	<b>Peer Reviewer</b>	
	• Conference on Neural Information Processing Systems (NeurIPS)	2021-2025
	• International Conference on Machine Learning (ICML)	2024-2025
	• International Conference on Artificial Intelligence and Statistics (AISTATS)	2024

<b>TEACHING</b>	<b>Teaching Assistant</b>	Cornell University
	• Learning, Inference, and Decision Making from Data	Spring 2022
	• Applied Machine Learning	Fall 2021
	<b>Teaching Fellow</b>	Harvard University
	• Mechanics and Special Relativity	Fall 2014
	• Data Science	Fall 2014
	• Linear Algebra and Real Analysis	Spring 2013
	• Algebra I	Fall 2013

**PUBLICATIONS FULL LIST** † - equal contribution, ‡ - alphabetical authors

Latest publications available on [Google Scholar](#).

## PREPRINTS

- [1] **Miruna Oprescu**, Brian M Cho, and Nathan Kallus. Efficient adaptive experimentation with non-compliance. *arXiv preprint arXiv:2505.17468*, 2025.
- [2] **Miruna Oprescu**, David K Park, Xihaier Luo, Shinjae Yoo, and Nathan Kallus. Gst-unet: Spatiotemporal causal inference with time-varying confounders. *arXiv preprint arXiv:2502.05295*, 2025.

## CONFERENCE PUBLICATIONS

- [1] **Miruna Oprescu** and Nathan Kallus. Estimating heterogeneous treatment effects by combining weak instruments and observational data. *Advances in Neural Information Processing Systems*, 37:118777–118806, 2025.

- [2] Andrew Bennett<sup>‡</sup>, Nathan Kallus<sup>‡</sup>, **Miruna Oprescu<sup>‡</sup>**, Wen Sun<sup>‡</sup>, and Kaiwen Wang<sup>‡</sup>. Efficient and sharp off-policy evaluation in robust markov decision processes. *Advances in Neural Information Processing Systems*, 37:112962–113000, 2025.
- [3] Andrew Bennett<sup>‡</sup>, Nathan Kallus<sup>‡</sup>, and **Miruna Oprescu<sup>‡</sup>**. Low-rank mdps with continuous action spaces. In *International Conference on Artificial Intelligence and Statistics*, pages 4069–4077. PMLR, 2024.
- [4] **Miruna Oprescu**, Jacob Dorn, Marah Ghoummaid, Andrew Jesson, Nathan Kallus, and Uri Shalit. B-learner: Quasi-oracle bounds on heterogeneous causal effects under hidden confounding. In *Proceedings of the 40th International Conference on Machine Learning*, pages 26599–26618. PMLR, 2023.
- [5] Nathan Kallus<sup>†</sup> and **Miruna Oprescu<sup>†</sup>**. Robust and agnostic learning of conditional distributional treatment effects. In *International Conference on Artificial Intelligence and Statistics*, pages 6037–6060. PMLR, 2023.
- [6] Soukayna Mouatadid, Paulo Orenstein, Genevieve Flaspohler, Judah Cohen, **Miruna Oprescu**, Ernest Fraenkel, and Lester Mackey. Adaptive bias correction for improved subseasonal forecasting. *Nature Communications*, 14(1):3482, 2023.
- [7] Keith Battocchi<sup>‡</sup>, Eleanor Dillon<sup>‡</sup>, Maggie Hei<sup>‡</sup>, Greg Lewis<sup>‡</sup>, **Miruna Oprescu<sup>‡</sup>**, and Vasilis Syrgkanis<sup>‡</sup>. Estimating the long-term effects of novel treatments. *Advances in Neural Information Processing Systems*, 34:2925–2935, 2021.
- [8] Genevieve E Flaspohler, Francesco Orabona, Judah Cohen, Soukayna Mouatadid, **Miruna Oprescu**, Paulo Orenstein, and Lester Mackey. Online learning with optimism and delay. In *International Conference on Machine Learning*, pages 3363–3373. PMLR, 2021.
- [9] **Miruna Oprescu<sup>†</sup>**, Vasilis Syrgkanis<sup>†</sup>, and Zhiwei Steven Wu<sup>†</sup>. Orthogonal random forest for causal inference. In *International Conference on Machine Learning*, pages 4932–4941. PMLR, 2019.
- [10] Vasilis Syrgkanis, Victor Lei, **Miruna Oprescu**, Maggie Hei, Keith Battocchi, and Greg Lewis. Machine learning estimation of heterogeneous treatment effects with instruments. In *Advances in Neural Information Processing Systems*, pages 15193–15202, 2019. **Spotlight presentation.**
- [11] **Miruna Oprescu<sup>†</sup>**, Vasilis Syrgkanis<sup>†</sup>, Keith Battocchi<sup>†</sup>, Maggie Hei<sup>†</sup>, and Greg Lewis<sup>†</sup>. EconML: A Machine Learning Library for Estimating Heterogeneous Treatment Effects. In *CausalML Workshop, NeurIPS*, 2019. **Spotlight presentation.**
- [12] K Arbour, **M Oprescu**, J Hakim, H Rizvi, M Leiserson, M Ginsburg, A Plodkowski, J Sauter, I Preeshagul, S Gillett, et al. Multifactorial Model to Predict Response to PD-(L) 1 Blockade in Patients with High PD-L1 Metastatic Non-Small Cell Lung Cancer. *Journal of Thoracic Oncology*, 14(10):S290, 2019.